

COMPTE RENDU TP3

MARIEM EJIWEN 23018

9 MAI 2025

EXO1 :

question1:

```
from pyspark.sql import SparkSession

spark = SparkSession.builder.appName("WorldBank")\
    .config("spark.jars.packages",
"org.apache.spark:spark-sql-kafka-0-10_2.12:3.4.1") \
    .getOrCreate()
```

Question 2 : Charger le fichier CSV dans un DataFrame, le convertir en RDD **country_data**, puis afficher les 5 premières lignes

```
df = spark.read.csv("world_bank_dataset.csv", header=True,
inferSchema=True)

country_data = df.rdd

country_data.take(5)

for row in rows:

    print(row)
```

```

org.slf4j#slf4j-api;2.0.6 from central in [default]
org.xerial.snappy#snappy-java;1.1.10.1 from central in [default]
-----
|               |      modules      |      artifacts      |
|      conf      | number| search|dwnlded|evicted|| number|dwnlded|
|-----|
|      default   |    11 |    0 |    0 |    0 ||    11 |    0 |
|-----|

:: retrieving :: org.apache.spark#spark-submit-parent-fee8671e-e857-4df9-bae9-db3c4f9cb9ae9
confs: [default]
0 artifacts copied, 11 already retrieved (0kB/26ms)
25/05/09 08:24:17 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
Row(Country='Brazil', Year=2010, GDP_USD=149322000000.0, Population=829020000.0, Life_Expectancy=66.7, Unemployment_Rate=3.81, CO2_Emissions=10.79, Access_to_Electricity_percent=76.76)
Row(Country='Japan', Year=2011, GDP_USD=1756270000000.0, Population=897010000.0, Life_Expectancy=61.4, Unemployment_Rate=17.98, CO2_Emissions=15.67, Access_to_Electricity_percent=67.86)
Row(Country='India', Year=2012, GDP_USD=1642688000000.002, Population=669850000.0, Life_Expectancy=69.1, Unemployment_Rate=16.02, CO2_Emissions=2.08, Access_to_Electricity_percent=81.08)
Row(Country='Mexico', Year=2013, GDP_USD=1189001000000.0, Population=113800000.0, Life_Expectancy=80.1, Unemployment_Rate=6.26, CO2_Emissions=19.13, Access_to_Electricity_percent=53.46)
Row(Country='India', Year=2014, GDP_USD=267302000000.0, Population=29710000.0, Life_Expectancy=62.7, Unemployment_Rate=3.1, CO2_Emissions=15.66, Access_to_Electricity_percent=82.17)
PS C:\Users\DELL\Desktop\bd-env\bd-env\work>

```

Question 3 : Afficher le nombre total d'enregistrements dans le RDD

```

print("nomnre",country_data.count())


```

```

org.lz4#lz4-java;1.8.0 from central in [default]
org.slf4j#slf4j-api;2.0.6 from central in [default]
org.xerial.snappy#snappy-java;1.1.10.1 from central in [default]
-----
|               |      modules      |      artifacts      |
|      conf      | number| search|dwnlded|evicted|| number|dwnlded|
|-----|
|      default   |    11 |    0 |    0 |    0 ||    11 |    0 |
|-----|

:: retrieving :: org.apache.spark#spark-submit-parent-98b05ad4-0b7f-4c0a-bdb6-36d25b0974cf
confs: [default]
0 artifacts copied, 11 already retrieved (0kB/27ms)
25/05/09 08:27:26 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform...
here applicable
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use
nomnre 200
PS C:\Users\DELL\Desktop\bd-env\bd-env\work>

```


Command failed: npm i --package

Question 4 : Quelle est la population et l'espérance de vie du Canada en 2014 ?

```
df.filter((df["Country"] == "Canada") & (df["Year"] == 2014)) \
    .select("Population", "Life_Expectancy") \
    .show()
```

```
+-----+-----+
|Population|Life_Expectancy|
+-----+-----+
|  4.1468E8|          56.4|
+-----+-----+
```

```
PS C:\Users\DELL\Desktop\bd-env\bd-env\work>
```

Question 5 : Afficher le nom de tous les pays dont le PIB dépasse 10 000 000 000 000 USD en 2015

```
df.filter((df["Year"] == 2015) & (df["GDP_USD"] > 1e13)) \
    .select("Country") \
    .distinct() \
    .show()
```

```

+-----+
|      Country      |
+-----+
|      Argentina    |
|           India    |
|         Nigeria    |
|           Italy     |
|        Indonesia   |
|   United Kingdom   |
+-----+
PS C:\Users\DELL\Desktop\bd-env\bd-env\work>

```

Question 6 : Trouver le nombre de pays dont le PIB dépasse 10 000 000 000 000 USD en 2015

```

print("pyas of pib>10 000 000 000 000 USD",df.filter((df["Year"] == 2015)
& (df["GDP_USD"] > 1e13)) \

.select("Country") \

.distinct() \

.count())

```

```

pyas of pib>10 000 000 000 000 USD 6
PS C:\Users\DELL\Desktop\bd-env\bd-env\work>

```

Question 7 : Afficher les 5 pays ayant le plus grand taux de chômage et leur taux

```
df.select("Country", "Unemployment_Rate") \  
  .orderBy(df["Unemployment_Rate"].desc()) \  
  .dropna(subset=["Unemployment_Rate"]) \  
  .distinct() \  
  .show(5)
```

```
+-----+-----+  
|   Country|Unemployment_Rate|  
+-----+-----+  
|Saudi Arabia|      2.62|  
|   Mexico|      16.86|  
|Australia|      3.12|  
|   Canada|      17.94|  
|Argentina|      18.52|  
+-----+-----+  
only showing top 5 rows
```

Question 8 : Quel est le PIB moyen des pays en 2014 ?

```
df.filter(df["Year"] == 2014) \  
  .select("GDP_USD") \  
  .groupBy().avg() \  
  .show()
```

```
+-----+
|avg(GDP_USD)|
+-----+
| 8.951888E12|
+-----+
```

EXO2:

Question 1 :

Charger les données JSON du fichier **reviews_nd.json** dans un RDD et afficher les 2 premiers documents.

```
import json

from pyspark.sql import SparkSession

spark = SparkSession.builder.appName("ReviewsRDD") \
    .config("spark.jars.packages",
"org.apache.spark:spark-sql-kafka-0-10_2.12:3.4.1") \
    .getOrCreate()

rdd_raw = spark.sparkContext.textFile("reviews_nd.json")

reviews_rdd = rdd_raw.map(lambda line: json.loads(line))

for review in reviews_rdd.take(2):
    print(review)
```

```
-----
|               | modules          | artifacts        |
|      conf     | number | search|dwnlded|evicted|| number|dwnlded|
|-----|-----|-----|-----|-----|-----|
|      default  |    11  |    0  |    0  |    0  ||    11  |    0  |
|-----|-----|-----|-----|-----|-----|

:: retrieving :: org.apache.spark#spark-submit-parent-e256020f-6c38-4330-b746-dc6b4c7d13dc
  confs: [default]
    0 artifacts copied, 11 already retrieved (0kB/19ms)
25/05/09 08:40:54 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
{'id_review': 1, 'film_id': 26, 'titre': 'The Platform', 'user_id': 332, 'user_name': 'Miranda Terry', 'user_age': 42, 'genre': 'Horreur', 'note': 8, 'date_production': '2020-09-11', 'date_review': '2022-06-30'}
{'id_review': 2, 'film_id': 34, 'titre': 'Free Guy', 'user_id': 84, 'user_name': 'Robin Thompson', 'user_age': 42, 'genre': 'Comédie', 'note': 8, 'date_production': '2021-08-13', 'date_review': '2022-11-11'}
PS C:\Users\DELL\Desktop\bd-env\bd-env\work>
```

Question 2 :

Calculer le nombre de films notés chaque année, ordonné par année croissante.

Cela se base sur le champ **"date_review"** . Il faut extraire l'année, grouper, puis compter.

```
from pyspark.sql import functions as F

df= spark.read.json("reviews_nd.json")

df.withColumn("year_review", F.year("date_review")) \

    .groupBy("year_review") \

    .count() \

    .orderBy("year_review") \

    .show()
```


year_review	count
2019	21
2020	171
2021	294
2022	474
2023	777
2024	263

Q3. Nombre de reviews et moyenne des notes par film

```
df.groupBy("film_id", "titre") \
    .agg(F.count("*").alias("nb_reviews"),
F.avg("note").alias("moyenne_note")) \
    .orderBy(F.col("nb_reviews").desc(), F.col("moyenne_note").desc()) \
    .show(5)
```

film_id	titre	nb_reviews	moyenne_note
60	Star Wars: The Ri...	33	5.484848484848484
45	Ghostbusters: Aft...	30	5.3
32	Luca	29	5.620689655172414
67	Jurassic World: D...	28	6.428571428571429
61	The Dark Knight	28	5.785714285714286

only showing top 5 rows

Q4. Moyenne d'âge des utilisateurs du film le plus noté

```

top_film=df.groupBy("film_id").count().orderBy(F.desc("count")).first()["film_id"]

df.filter(df["film_id"] == top_film) \

    .agg(F.avg("user_age").alias("moyenne_age")) \

    .show()

```

```

+-----+
|      moyenne_age      |
+-----+
|44.03030303030303|
+-----+

```

Q5. Statistiques par utilisateur (top 4)

```

df.groupBy("user_id", "user_name") \

    .agg(

        F.max("note").alias("max_note"),

        F.min("note").alias("min_note"),

        F.avg("note").alias("moyenne_note")

    ) \

    .orderBy(F.col("moyenne_note").desc()) \

    .show(4)

```

```

+-----+-----+-----+-----+-----+
|user_id|  user_name|max_note|min_note|moyenne_note|
+-----+-----+-----+-----+-----+
|   172| Sarah Lucero|    10|     9|        9.5|
|   150|Melissa Romero|    10|     9|        9.5|
|   124|  Sheena Hall|    10|     9|        9.5|
|   442| Samuel Frank|    10|     9|       9.25|
+-----+-----+-----+-----+-----+
only showing top 4 rows

```

Q6. Mois avec le plus de reviews

```

df.withColumn("mois", F.month("date_review")) \
  .groupBy("mois") \
  .count() \
  .orderBy(F.col("count").desc()) \
  .show(1)

```

```

+----+-----+
|mois|count|
+----+-----+
|   1|  249|
+----+-----+
only showing top 1 row

```

Q7. Genre le plus populaire (en nombre de reviews)

```

df.groupBy("genre") \
  .count() \
  .orderBy(F.desc("count")) \

```

```
.show(1)
```

```
+-----+-----+
| genre|count|
+-----+-----+
|Action|  662|
+-----+-----+
only showing top 1 row
```

Q8. Genre le mieux noté (en moyenne)

```
df.groupBy("genre") \
  .agg(F.avg("note").alias("moyenne_note")) \
  .orderBy(F.desc("moyenne_note")) \
  .show(1)
```

```
+-----+-----+
|      genre|moyenne_note|
+-----+-----+
|Fantastique|         6.72|
+-----+-----+
only showing top 1 row
```

Q9. Pour chaque année de production, film ayant reçu le plus de notes

```
from pyspark.sql.window import Window

window = Window.partitionBy("date_production").orderBy(F.desc("count"))
```

```
df.groupBy("date_production", "film_id", "titre") \

    .count() \

    .withColumn("rang", F.row_number().over(window)) \

    .filter(F.col("rang") == 1) \

    .orderBy("date_production") \

    .show()
```

```
+-----+-----+-----+-----+-----+
|date_production|film_id|      titre|count|rang|
+-----+-----+-----+-----+-----+
| 2019-05-03|    88|Black Panther: Th...|   10|  1|
| 2019-05-05|    91|Guardians of the ...|   25|  1|
| 2019-06-15|    85|Jurassic World: T...|   22|  1|
| 2019-06-30|    94|Pixar's Magical A...|   21|  1|
| 2019-07-14|    92|Indiana Jones: Th...|   20|  1|
| 2019-07-28|    99|The Adventures of...|   23|  1|
| 2019-08-17|    84|Justice League: U...|   18|  1|
| 2019-10-05|    87|Spider-Man: The W...|   23|  1|
| 2019-10-06|    97|DC Universe: Lege...|   10|  1|
| 2019-11-02|    90|Thor: War of the ...|   15|  1|
| 2019-11-16|    93|Fantastic Beasts:...|   19|  1|
| 2019-11-17|    86|The Flash: Beyond...|   17|  1|
| 2019-11-22|    96|Disney's Enchante...|   19|  1|
| 2019-12-01|   100|DreamWorks Animat...|   18|  1|
| 2019-12-22|    98|Minions: The Ques...|   28|  1|
| 2019-12-25|    83|Avatar 3: The Fin...|   21|  1|
| 2020-01-08|    19|      Underwater|   13|  1|
| 2020-01-17|     1|      Dolittle|   15|  1|
| 2020-01-24|     7|    The Gentlemen|   22|  1|
| 2020-01-25|     5|    Birds of Prey|   21|  1|
+-----+-----+-----+-----+-----+
only showing top 20 rows
```

EXO3:

Question 1 : Qui est l'employé le plus actif dans l'envoi des emails ?

```
df_enron.groupBy("Sender") \
    .count() \
    .orderBy(F.desc("count")) \
    .show(1)
```

```
+-----+-----+
|Sender|count|
+-----+-----+
|  NULL|19038|
+-----+-----+
only showing top 1 row
```

Q2. Nombre d'emails envoyés avant et après le 2 décembre 2001

```
before = df_enron.filter(F.col("Date") < "2001-12-02").count()
after = df_enron.filter(F.col("Date") >= "2001-12-02").count()

print(f"Avant le 02/12/2001 : {before} emails")
print(f"Après le 02/12/2001 : {after} emails")
```

```
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
Avant le 02/12/2001 : 334915 emails
Après le 02/12/2001 : 173440 emails
```

Q3. Heure d'activité la plus intense entre 8h et 17h

```
df = df_enron.withColumn("Date", to_timestamp("Date", "dd-MM-yyyy  
HH:mm:ss"))  
  
df.filter((F.col("Hour") >= 8) & (F.col("Hour") <= 17)) \  
  
  .groupBy("Hour") \  
  
  .count() \  
  
  .orderBy(F.desc("count")) \  
  
  .show(1)
```

```
+---+-----+  
|hour|count|  
+---+-----+  
|   8|43042|  
+---+-----+  
only showing top 1 row
```

Q4. Mentions de "fraud" et "bankruptcy" dans le sujet

```
fraud_count =  
df_enron.filter(F.lower(F.col("Subject")).contains("fraud")).count()  
  
bankruptcy_count =  
df_enron.filter(F.lower(F.col("Subject")).contains("bankruptcy")).count()  
  
print(f"'fraud' mentionné : {fraud_count} fois")
```

```
print(f'"bankruptcy" mentionné : {bankruptcy_count} fois')
```

```
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR,
'fraud' mentionné : 17 fois
'bankruptcy' mentionné : 691 fois
PS C:\Users\DELL\Desktop\bd-env\bd-env\work> 
```

Q5. Emails envoyés par **Jeff.Skilling@enron** avec le sujet "FREE LUNCH ON FRIDAY!"

```
from pyspark.sql.functions import lower, col

df_enron.filter(

    (lower(col("Sender")).like("%skilling%")) &

    (lower(col("Subject")) == "free lunch on friday!")

).select("Date", "Sender", "Subject").show(truncate=False)
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+
|09-08-2001 19:29:59|Ken Lay and Jeff Skilling@ENRON <IMCEANOTES-Ken+20Lay+20and+20Jeff+20Skilling+40ENRON@ENRON.com>|FREE LUNC
H ON FRIDAY!|
|09-08-2001 19:29:57|Ken Lay and Jeff Skilling@ENRON                                     |FREE LUNC
H ON FRIDAY!|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+
```

Q6. Ratio d'emails internes vs externes

```
df = df_enron.withColumn("is_internal",
F.col("Recipients").contains("@enron"))

nb_total = df.count()
```



```

nb_internal = df.filter(F.col("is_internal")).count()

nb_external = nb_total - nb_internal

ratio = nb_internal / nb_external if nb_external else "Infini"

print(f"Emails internes : {nb_internal}")

print(f"Emails externes : {nb_external}")

print(f"Ratio interne/externe : {ratio}")

```

```

Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
Emails internes : 49101
Emails externes : 459254
Ratio interne/externe : 0.10691469208760294

```

Exo4:

Q1. Transformer le dataset en RDD

```

import json

from pyspark.sql import SparkSession

from pyspark.sql import functions as F

from pyspark.sql.window import Window

from pyspark.sql.functions import to_timestamp, hour

from pyspark.sql.functions import lower, col

```

```

spark = SparkSession.builder.appName("AirbnbAnalysis")\
    .config("spark.jars.packages",
"org.apache.spark:spark-sql-kafka-0-10_2.12:3.4.1") \
    .getOrCreate()

df = spark.read.csv("airbnb.csv", header=True, inferSchema=True)

rdd = df.rdd

rdd.take(1)

print(rdd.take(1))

)

```

```

[Row(listing_id=2352, date=datetime.date(2023, 12, 21), available='f', price='$89.00', adjusted_price=None, minimum_nights=2, maximum_nights=1125)]

```

Q2. Quelle est la première date ? Le logement était-il disponible ?

```

first_date_row = df.orderBy("date").select("date", "available").first()

print(f"Première date : {first_date_row['date']}, Disponible : {first_date_row['available']}")

```

```

To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
Première date : 2023-12-20, Disponible : f
PS C:\Users\DELL\Desktop\bd-env\bd-env\work>

```

Q3. Combien d'annonces étaient disponibles le 31/12/2023 ?

```

print("resultat:",df.filter((F.col("date") == "2023-12-31") &
(F.col("available") == "t")).count())

```

```
resultat: 603
```

```
PS C:\Users\DELL\Desktop\bd-env\bd-env\work>
```

Q4. Annonces ayant `maximum_nights > 365`

```
df.filter(F.col("maximum_nights") > 365).select("listing_id",  
"maximum_nights").distinct().show(  
)
```

listing_id	maximum_nights
2701606	1125
688398	1125
8585748	1125
12674189	1125
4543524	1125
5793552	1125
8127197	1125
18341230	1125
18502751	1125
5214831	1125
13119454	1125
17174345	1124
20009913	1125
21052562	1125
21888524	1000
1155898	1125
7836643	1125
21052562	1125
21888524	1000
1155898	1125
7836643	1125
17615195	1125
4065776	730
6176433	1125

Q5. Pour `listing_id = 2352`, afficher les 10 premières dates non disponibles

```
df.filter((F.col("listing_id") == 2352) & (F.col("available") == "f")) \

.orderBy("date") \

.select("date") \

.show(10)
```

```
+-----+
|      date|
+-----+
|2023-12-21|
|2023-12-22|
|2023-12-23|
|2023-12-24|
|2023-12-25|
|2023-12-26|
|2023-12-27|
|2023-12-28|
|2023-12-29|
|2023-12-30|
+-----+
only showing top 10 rows
```

Q6. 10 annonces aléatoires où price == adjusted_price

```
from pyspark.sql.functions import rand

df.filter((F.col("price") == F.col("adjusted_price"))) \

.orderBy(rand()) \

.select("listing_id", "date", "price", "adjusted_price") \

.show(10, truncate=False)
```

```

+-----+-----+-----+-----+
|listing_id|date|price|adjusted_price|
+-----+-----+-----+-----+
|816624685831405709|2024-07-06|$1,200.00|$1,200.00|
|816624685831405709|2024-09-23|$1,200.00|$1,200.00|
|816614573653730356|2024-11-21|$1,000.00|$1,000.00|
|816624685831405709|2024-08-29|$1,200.00|$1,200.00|
|816614573653730356|2024-05-22|$1,000.00|$1,000.00|
|816624685831405709|2024-11-02|$1,200.00|$1,200.00|
|816624685831405709|2024-10-25|$1,200.00|$1,200.00|
|816624685831405709|2024-02-13|$1,200.00|$1,200.00|
|816624685831405709|2024-08-26|$1,200.00|$1,200.00|
|816624685831405709|2024-09-13|$1,200.00|$1,200.00|
+-----+-----+-----+-----+
only showing top 10 rows

```

Question 7 :Combien d’annonces ont une minimum_nights supérieure à la moyenne globale ?

```

from pyspark.sql import functions as F

moyenne_min_nights = df.select(F.avg("minimum_nights")).first()[0]

print("ressltat:",df.filter(F.col("minimum_nights") >
moyenne_min_nights).count())

```

```

Setting default log level to WARN .
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
ressltat: 333465
PS C:\Users\DELL\Desktop\bd-env\bd-env\work>

```